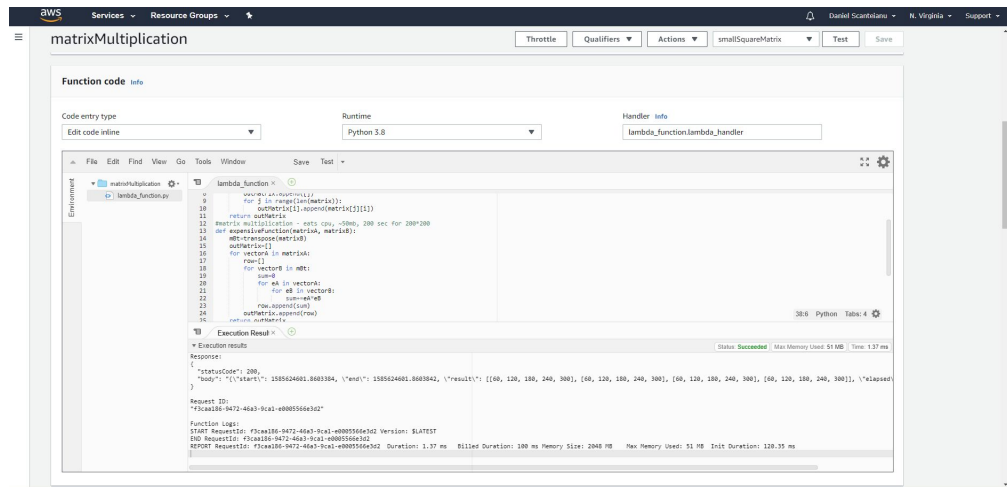


Black Box Benchmarking FAAS Platforms

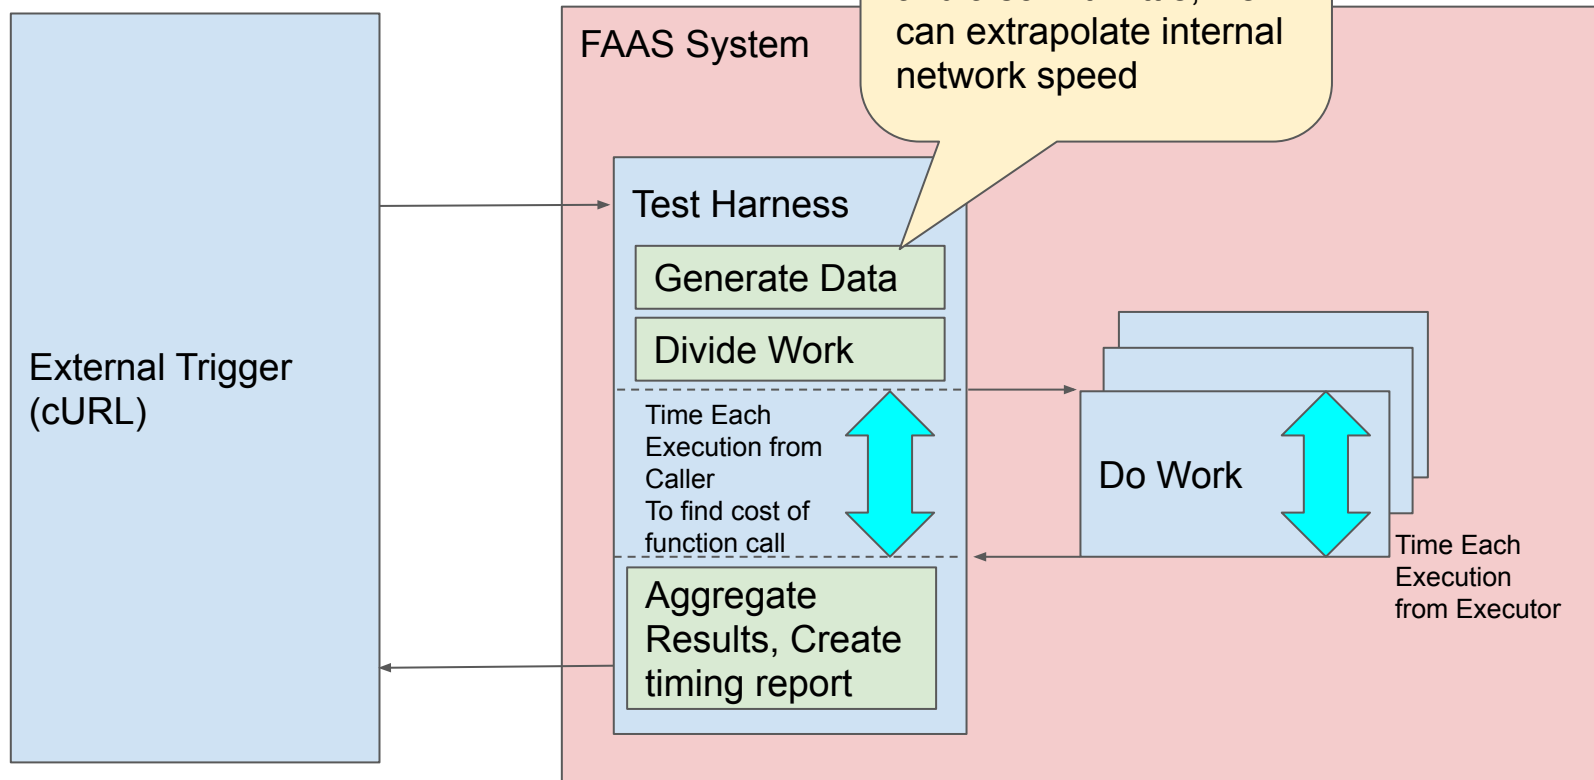
Daniel Scanteianu OS/2 Project

What is FAAS/Serverless?

- Function As A Service
- Serverless
 - Cloud provider handles server
 - You write some code
 - Cloud provider deploys
 - Function gets triggered
 - Cloud provider handles scaling up and scaling down
 - You get billed per invocation/execution time
- Uses containers under the hood
 - "Preheated" runtimes accept function run requests



Benchmark Design



Where we left off

- Goal: find out how different FAAS systems perform under different load patterns
 - Matrix Multiplication Test
 - Meant to find out who had a faster cpu
 - Gathered other metrics
 - AWS vs Google - we found AWS was faster in the function, but Google spent less time between the function and the harness
 - Add IBM - based on Microsoft results from other study, IBM is a more production-ready choice
 - Naive matrix multiplication has been pointed out to not be the most valid test

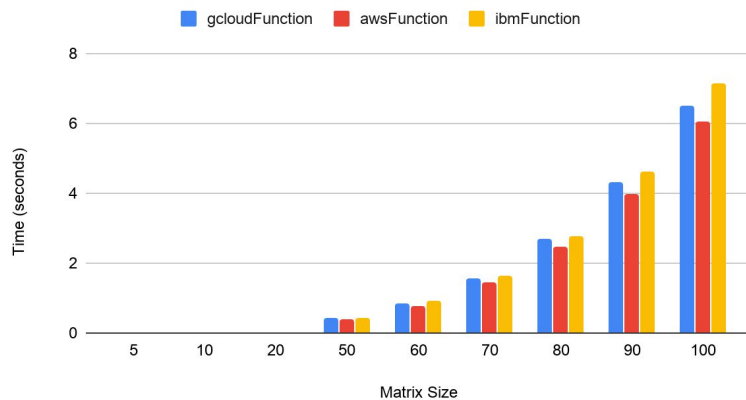
Other Work

- Reliability (failure rate) when running lots of requests:
 - AWS 0.954%, GCF 2.777%, IBM 2.775, MS 95.976%
 - https://www.ise.tu-berlin.de/fileadmin/fg308/publications/2020/Online_Preprint_SAC_2020_Benchmarking_Elasticity_of_FaaS_Platforms_as_a_Foundation_for_Objective_driven_Design_of_Serverless_Applications.pdf
- Startup delays for google cloud vs aws - AWS<Google<IBM
 - <https://oaciss.uoregon.edu/icpp18/views/includes/files/pos115s2-file1.pdf>
- Evaluation of AWS, Google, MSFT on cpu intensive tasks (video processing, ml) shows that for the most part outperforms others, especially toward high end
 - <http://jckim.me/assets/paper/FunctionBench%20-%20A%20Suite%20of%20Workloads%20for%20Serverless%20Cloud%20Function%20Service.pdf>
- Elasticity - ability to serve requests at same speed as load increases
<https://www.cs.colostate.edu/~shrdeep/papers/ServerlessComputing-IC2E-2018.pdf> (also other tests)
- Total time vs execution time (USENIX)! In aws step functions
 - <https://www.usenix.org/system/files/conference/atc18/atc18-akkus.pdf>

Test: Matrix Multiplication

- Naive Matrix Multiplication Algorithm
 - 2gb instances, $O(n^2) \cdot n^2$ cells in matrix
- Headline results
 - AWS has a faster CPU than Google or IBM
 - Statistically significant difference
 - ~18% faster - explainable by slightly newer CPUs
 - Presumably very similar underlying systems
 - Enterprise grade vms running containers which host functions, x86 based, linux
- Other observations were more interesting

Time Measured In Function

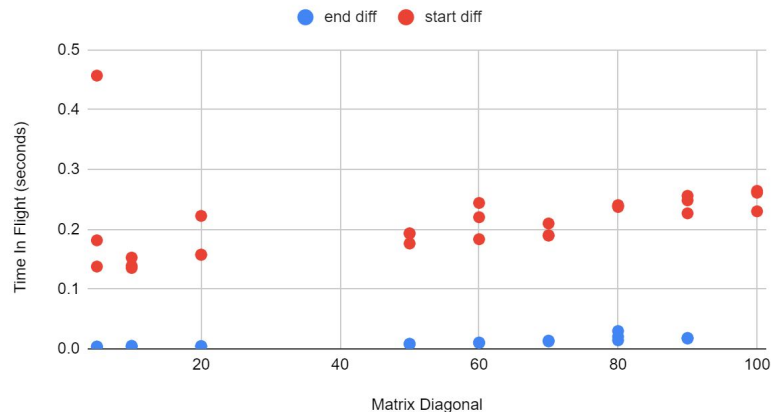


aws < gc	aws < ibm	gc < ibm
p=0.0001434475609	p=0.0003607757513	p=0.001313550811


Observations

- Measured start/end harness vs function time
- Noticed aws end time was pretty low, but mild upward trend as the matrix size increased
- Noticed that there was a bigger spread in Google timings than in Amazon timings
- New experiment - which provider has fastest (and most consistent) network?

Start/End Processing Time Vs Matrix Size



Start/End Processing Time Vs Matrix Size




Test: Large Network

- Very simple test
 - Generate huge list
 - Send over the wire to function
 - Return length of list
 - Assuming JSON serialization, estimate network
- Isolates network component
- Desired output
 - Learn about the network
 - Can estimate size of packet sent over the wire
 - Find out both approximate network speed
 - Includes ser/de because it's a list encoded in json
 - "Real" cost of passing data to a function
 - Learn about network variability

Google Cloud vs AWS Time In Network

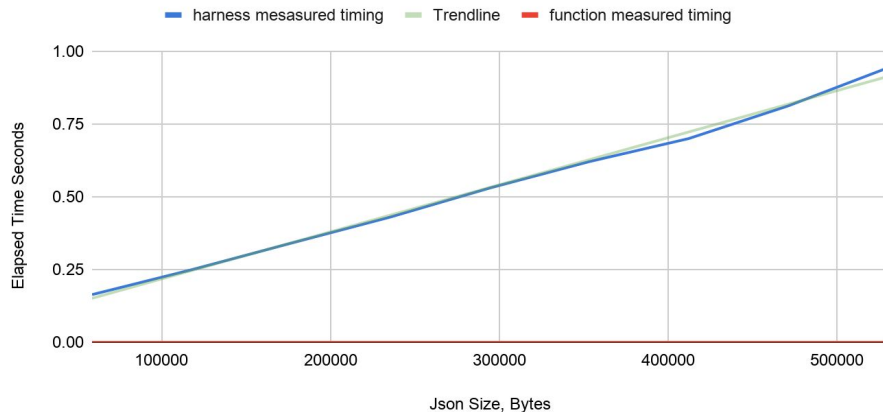
Difference Between Function and Harness Time




Test: Large Network

- AWS is surprisingly low in variability
- Last 3 Google data points might indicate that they let larger data have more bandwidth

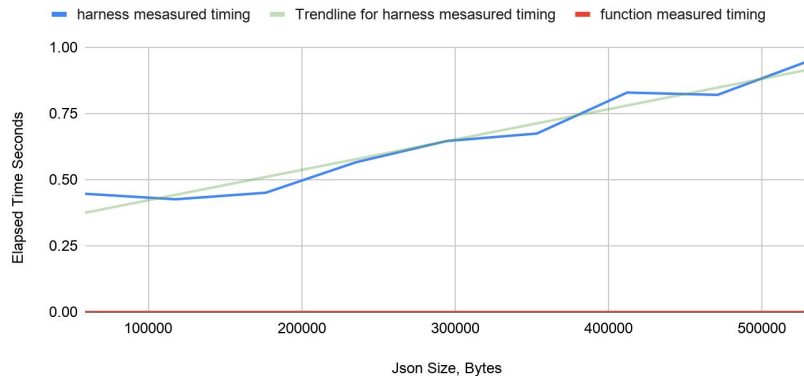
AWS: Execution Time vs List Size



Gcloud: Execution Time vs List Size



IBM: Execution Time vs List Size



Test: Large Network

- Google is fastest
- AWS - 99.25% of variance in response timing due to packet size
 - Surprising - means network is super super consistent
 - Throttling or underused?
- Extremely unlikely difference in means is due to random variance


Provider	Network Speed (MBPS)	Startup Latency	R^2
AWS	4.910699466	0.05582174593	0.9925222419
Google	7.025047395	0.04267904891	0.8691661488
IBM	5.352188787	0.3073757185	0.7667496019

aws < gc	aws < ibm	gc > ibm
p=0.0	p<0.00001	p=0.0

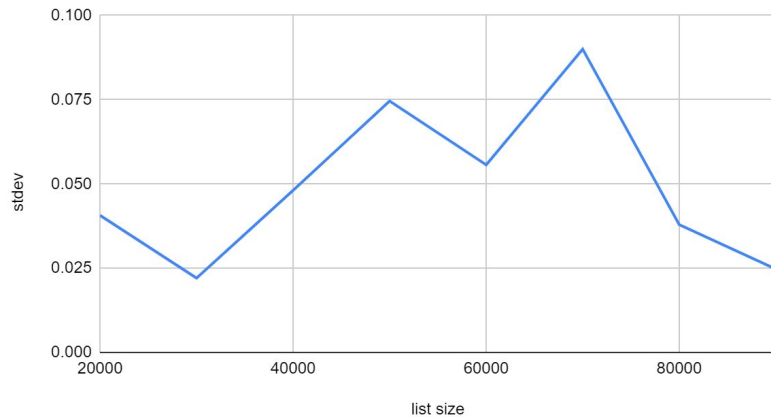
Test: Large Network

- AWS looks like I expect
 - More data = more opportunity for network to vary
- Google very consistent for bigger requests
- Initial startup latency clouding measurement


aws list size vs stdev



gcloud list size vs stdev



ibm list size vs stdev



Conclusions: Large Network

- Amazon is very predictable and consistent
 - Slowest network, but seems that the only factor affecting start up time increase was increase in data size
 - Boto3 library as opposed to raw http request
 - Internal throttling may be happening
- Google is faster than other two, and might have benefits for larger data sets
- Was able to get actual MBPS figure for all 3 providers
- All 3 tests included built in lambda ser/de time
 - No really good way to bypass this

Perspective on FAAS platforms

- Small, distinct learning curve for each individual provider
- AWS has a much larger selection of runtimes available than other providers
 - More production-ready system than other providers
 - Few failures
 - Responsive UI
 - Super consistent
- Google is powerful, but rougher around the edges
 - Glitchy UI was hard to work with
- Pick one ecosystem and switch with it
 - Migrating functions is bad
 - Migrating communication/IAMs/Databases is worse

Conclusions on FAAS as a tool

- Very easy to learn and spin up services
- All IBM experiments fit within free tier limits
 - Generous offering, especially for students and others with limited resources
- Required much less time spent in configuring a web server
 - Very easy to add rest endpoint exposed in Google or IBM directly to an otherwise static website
- Automatic scaling is a very attractive offer
 - Huge devotion of engineer time to scalability and scaling in industry, and very complex planning is involved
- Potential for vendor lock in
- Well suited for people building applications
 - Use the right language for the function, common interface
- Probably way more reliable than having/maintaining your own servers

Future

- Case for implementing managed FAAS in industry
 - Multiple teams who own microservices might benefit
- I might try to make some contribution to some open source FAAS platform over the summer
- Any Questions?